


Harness-Level Latency, Tokens, and Estimated Cost Across 15 Cases



Latency means exclude Original PG Agent on DCI and Claude Code on DocVQA; token means exclude only the zero-token DocVQA runtime failure.
Cost uses DeepSeek V4 Pro list prices checked 2026-07-15: \$0.435/M uncached input, \$0.003625/M cache read, and \$0.87/M output; actual provider bills may differ.